

**Name- Dhyey Findoriya**  
**Roll No.- B22EE024**

**Comp Arch Lab-1**

---

**Machine-1 Specifications:**  
**OS-Fedora Linux 39**

Architecture:	x86_64
CPU op-mode(s):	32-bit, 64-bit
Address sizes:	39 bits physical, 48 bits virtual
Byte Order:	Little Endian
CPU(s):	8
On-line CPU(s) list:	0-7
Vendor ID:	GenuineIntel
Model name:	11th Gen Intel(R) Core(TM) i5-1135G7 @ 2.40GHz
CPU family:	6
Model:	140
Thread(s) per core:	2
Core(s) per socket:	4
Socket(s):	1
Stepping:	1
CPU(s) scaling MHz:	24%
CPU max MHz:	4200.0000
CPU min MHz:	400.0000
BogoMIPS:	4838.40

**About the code:**

**MandelBrot Problem(From GFG)**

- The Mandelbrot set is a collection of complex numbers where the iterative sequence  $z_{n+1} = z_n^2 + c$  remains bounded.
- Each point represents a complex number  $c = x + yi$ , with the set visualized as a fractal image.
- The function is iteratively applied to determine if the sequence diverges for each point.
- Points are colored based on the number of iterations before divergence, creating a detailed fractal.
- Uses complex numbers, iterative functions, and the concept of fractals to define and visualize the set.

### Analyzing the metrics:

Metric	Value	Details
Task Clock	657.10ms	
CPUs Utilized	0.930	
Context Switches	50	76.092 /sec
CPU Migrations	6	9.131 /sec
Page Faults	76	115.660 /sec
Cycles	2,236,805,514	3.404 GHz
Instructions	4,406,708,556	1.97 instructions per cycle
Branches	617,164,238	939.226 M/sec
Branch Misses	25,901,643	4.20% of all branches
TopdownL1		
- Backend Bound	13.1%	
- Bad Speculation	24.8%	
- Frontend Bound	22.2%	
- Retiring	40.0%	
Time Elapsed	1.006904578 seconds	
User Time	0.646168000 seconds	
System Time	0.006917000 seconds	

### Analyzing using linux command:

Metric	cd	pwd	ls
Task Clock	2.03 msec	1.22 msec	0.67 msec
CPUs Utilized	0.768	0.203	0.327
Context Switches	0	0	2
CPU Migrations	0	0	0
Page Faults	138	69	97
Cycles	5,905,836	2,674,957	2,488,525
Instructions	2,885,653	1,211,022	2,110,343
Branches	591,225	253,999	429,366
Branch Misses	15,941	7,953	14,165
TopdownL1			
- Backend Bound	39.6%	48.6%	32.5%
- Bad Speculation	5.6%	5.5%	10.7%
- Frontend Bound	44.7%	35.6%	37.5%
- Retiring	10.2%	10.2%	19.2%
Time Elapsed	0.002646429 seconds	0.006021428 seconds	0.002059844 seconds
User Time	0.001372000 seconds	0.000000000 seconds	0.000000000 seconds
System Time	0.001379000 seconds	0.003443000 seconds	0.001275000 seconds

## Analysis

### 1) Resources Utilization:

- **cd**: Shows the highest task clock time (2.03 msec) and CPU cycles (5,905,836), indicating more resource usage compared to **pwd** and **ls**. It has the lowest user and system time, but higher page faults.
- **pwd**: Consumes less task clock time (1.22 msec) and CPU cycles (2,674,957) than **cd**. It also has a lower number of page faults and branch misses.
- **ls**: Has the lowest task clock time (0.67 msec) and a relatively high number of CPU cycles (2,488,525) and instructions (2,110,343). It has a moderate number of context switches and page faults compared to **pwd** and **cd**.

### 2) Efficiency

- **ls** appears to be the most efficient in terms of task clock and system time. Despite having a moderate number of context switches and page faults, it uses fewer resources compared to **cd**.
- **pwd** has low task clock time and system time, but the number of page faults and branch misses are relatively higher than **ls**.
- **cd** has the highest task clock time and CPU cycles, indicating it might be doing more work or involving more complex operations.

### 3) Performance Characteristics

- **ls**: Likely involves scanning directory contents, which explains the higher cycles and instructions. Despite this, it is efficient in terms of time elapsed.
- **pwd**: Primarily involves retrieving and printing the current directory path, resulting in lower resource usage and faster execution.
- **cd**: Although it appears simple, it might involve filesystem operations or other underlying operations that contribute to its higher resource usage.

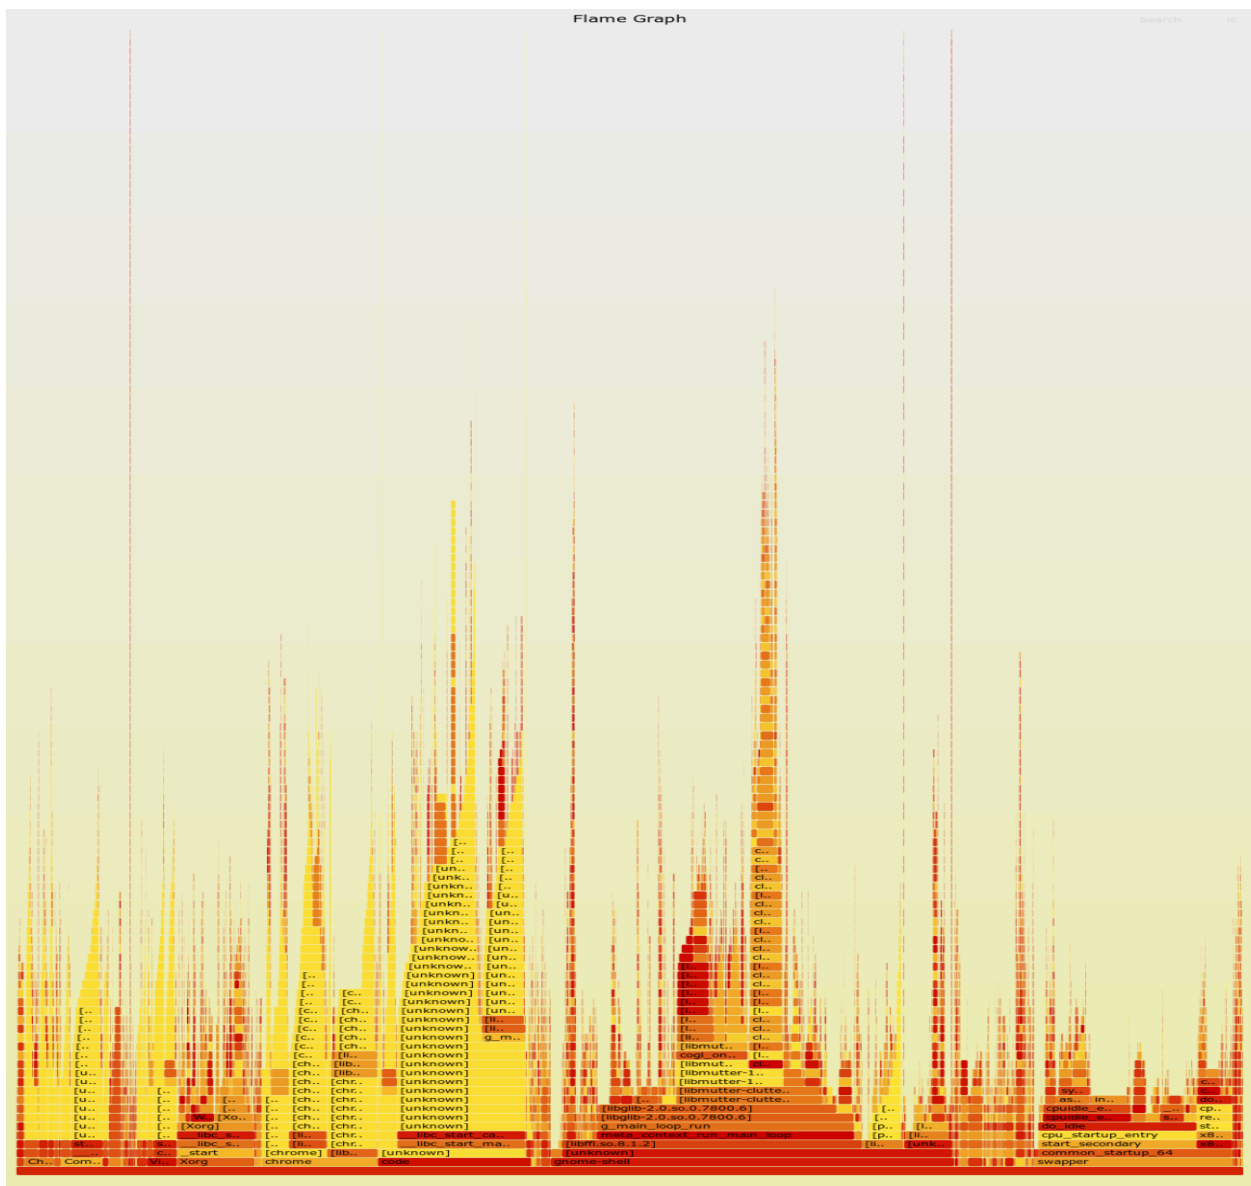
### 4) Potential Bottlenecks

- **cd**: The high number of CPU cycles and page faults suggest potential inefficiencies or additional overhead associated with changing directories.
- **pwd**: Higher branch misses and page faults could be an area for further investigation to improve efficiency.

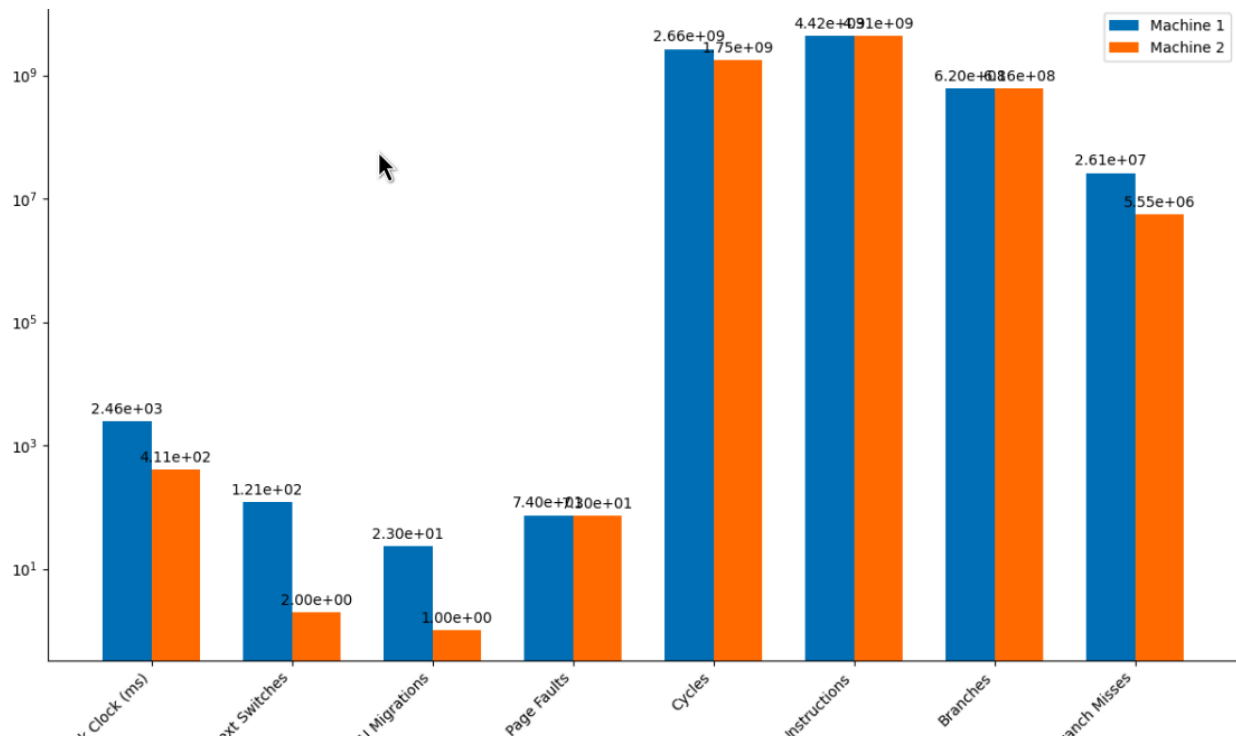
## My Observations

- **For `cd`:** Investigate the filesystem operations or additional tasks performed during directory changes. Optimization might involve reducing overhead or improving efficiency in handling filesystem operations.
- **For `pwd`:** While efficient, reducing branch misses and page faults can further optimize performance, especially for frequent use.

## Flamegraph



## Visualizing the data:



## Perf Bench command:

`Perf bench syscall all`

Benchmark	Calls Executed	Total Time	Time per Operation	Operations per Second
<code>syscall/basic</code>	10,000,000	1.374 sec	0.137433 $\mu$ s/op	7,276,299 ops/sec
<code>syscall/getpgid</code>	10,000,000	1.453 sec	0.145394 $\mu$ s/op	6,877,853 ops/sec
<code>syscall/fork</code>	10,000	12.253 sec	1225.313300 $\mu$ s/op	816 ops/sec
<code>syscall/execve</code>	10,000	11.613 sec	1161.368400 $\mu$ s/op	861 ops/sec

- **Calls Executed:** Number of system calls performed in the benchmark.
- **Total Time:** Total time taken to execute the benchmark.
- **Time per Operation:** Average time taken for each individual system call.
- **Operations per Second:** Number of system calls handled per second.

## Machine-2 Specifications: OS-Kali Linux

Architecture:	x86_64
CPU op-mode(s):	32-bit, 64-bit
Address sizes:	48 bits physical, 48 bits virtual
Byte Order:	Little Endian
CPU(s):	12
On-line CPU(s) list:	0-11
Vendor ID:	AuthenticAMD
Model name:	AMD Ryzen 5 5625U with Radeon Graphics
CPU family:	25
Model:	80
Thread(s) per core:	2
Core(s) per socket:	6
Socket(s):	1
Stepping:	0
CPU(s) scaling MHz:	31%
CPU max MHz:	4388.0000
CPU min MHz:	400.0000
BogoMIPS:	4591.50

### Analyzing the metrics:

Metrics	Value	Details
Task Clock	411.07 msec	0.997 CPUs utilized
Context Switches	2	4.865 /sec
CPU Migrations	1	2.433 /sec
Page Faults	73	177.585 /sec
Cycles	1,746,945,076	4.250 GHz
Stalled Cycles (Frontend)	35,530,577	2.03% frontend cycles idle
Instructions	4,312,074,188	2.47 insn per cycle
Stalled Cycles per Instruction	0.01	
Branches	616,022,963	1.499 G/sec
Branch Misses	5,550,725	0.90% of all branches
Elapsed Time	0.412235896 sec	
User Time	0.408745000 sec	
System Time	0.004007000 sec	

### Observation

	Machine-1(Fedora)	Machine-2(Kali)
Run time	1.0069sec	0.4122sec



## 1) Hardware Differences:

- **Kali Linux Machine:** 12 CPU cores (6 cores, 12 threads) with a maximum clock speed of 4.39 GHz.
- **Fedora Linux Machine:** 8 CPU cores (4 cores, 8 threads) with a maximum clock speed of 4.20 GHz.
- The Kali machine has more cores and threads, which can lead to better parallel processing capabilities, especially in multi-threaded applications, reducing execution time.

## 2) CPU Architecture:

- **Kali:** AMD Ryzen 5 5625U (Zen 3 architecture).
- **Fedora:** Intel Core i5-1135G7 (Tiger Lake architecture).
- The architectures of the CPUs can impact performance based on how well each is optimized for certain tasks. Ryzen processors are known for better multi-threaded performance, which might explain the faster execution on the Kali machine.

## 3) Kernel Optimization:

- **Kali Linux:** Often uses a custom kernel optimized for low latency and quick response times, which benefits short, intensive tasks. This can result in faster execution times for certain workloads, especially those that are CPU-bound.
  - **Fedora Linux:** Uses a more general-purpose kernel optimized for a balance between performance, stability, and security. This might not be as finely tuned for performance in high-demand scenarios, leading to slightly slower execution times.
- 
-